

## Assignment 11.1

### Contents

Transfer data between Mysql and HDFS (Import and Export) using Sqoop .....	1
Import: Transfer of data from Mysql to HDFS .....	1
Export: Transfer of data from HDFS to Mysql.....	4
Transfer data between Mysql and Hive (Import and Export only selected columns) using Sqoop.....	5
Import: Transfer of data from Mysql to hive. ....	5
Export: Transfer of data from hive to Mysql.....	9

### Transfer data between Mysql and HDFS (Import and Export) using Sqoop

#### Import: Transfer of data from Mysql to HDFS

##### Steps:

- **Create table in Mysql**  
create database assignment11;  
use assignment11;  
create table employee  
(id int, name varchar(20),  
age int, skill varchar(20),  
salary integer(10));

```
mysql> show databases;
+-----+
| Database |
+-----+
| information_schema |
| b1 |
| db1 |
| metastore |
| mysql |
| test |
+-----+
6 rows in set (0.00 sec)

mysql> create database assignment11;
Query OK, 1 row affected (0.00 sec)

mysql> use assignment11;
Database changed
mysql> create table employee
-> (id int, name varchar(20),
-> age int, skill varchar(20),
-> salary integer(10));
Query OK, 0 rows affected (0.01 sec)

mysql> show databases;
+-----+
| Database |
+-----+
| information_schema |
| assignment11 |
| b1 |
| db1 |
| metastore |
| mysql |
| test |
+-----+
7 rows in set (0.00 sec)
```

- **Insert Data to created table: 'employee'**

```
insert into employee values(1, 'Mohan',25,'Big Data & Hadoop',30000);
insert into employee values(2, 'Ramu',27,'AI',50000);
insert into employee values(3, 'Ravi',30,'Java',60000);
insert into employee values(4, 'Akshith',22,'Automation',35000);
insert into employee values(5, 'Shyam',35,'C',40000);
insert into employee values(6, 'Priya',28,'.Net',50000);
insert into employee values(7, 'Madhu',27,'DBA',70000);
```

**Grant all privileges to root user and commit the data.**

```
grant all on *.* to 'root'@'localhost' with grant option;
flush privileges;
commit;
exit;
```

```
mysql> select * from employee;
Empty set (0.00 sec)

mysql> insert into employee values(1, 'Mohan',25,'Big Data & Hadoop',30000);
Query OK, 1 row affected (0.00 sec)

mysql> insert into employee values(2, 'Ramu',27,'AI',50000);
Query OK, 1 row affected (0.00 sec)

mysql> insert into employee values(3, 'Ravi',30,'Java',60000);
Query OK, 1 row affected (0.00 sec)

mysql> insert into employee values(4, 'Akshith',22,'Automation',35000);
Query OK, 1 row affected (0.00 sec)

mysql> insert into employee values(5, 'Shyam',35,'C',40000);
Query OK, 1 row affected (0.00 sec)

mysql> insert into employee values(6, 'Priya',28,'.Net',50000);
Query OK, 1 row affected (0.00 sec)

mysql> insert into employee values(7, 'Madhu',27,'DBA',70000);
Query OK, 1 row affected (0.00 sec)

mysql> select * from employee;
+-----+-----+-----+-----+-----+
| id   | name   | age  | skill                | salary |
+-----+-----+-----+-----+-----+
| 1    | Mohan  | 25   | Big Data & Hadoop    | 30000  |
| 2    | Ramu   | 27   | AI                   | 50000  |
| 3    | Ravi   | 30   | Java                 | 60000  |
| 4    | Akshith| 22   | Automation           | 35000  |
| 5    | Shyam  | 35   | C                    | 40000  |
| 6    | Priya  | 28   | .Net                 | 50000  |
| 7    | Madhu  | 27   | DBA                  | 70000  |
+-----+-----+-----+-----+-----+
7 rows in set (0.00 sec)
```

- Use below query to import data from Mysql table 'employee' to HDFS directory 'sqoopout'

```
sqoop import --connect jdbc:mysql://localhost/assignment11 \
--username 'root' -P --table 'employee' --target-dir '/sqoopout' \
-m 1;
```

```
[acadgild@localhost ~]$ sqoop import --connect jdbc:mysql://localhost/assignment11 \
> --username 'root' -P --table 'employee' --target-dir '/sqoopout' \
> -m 1;
Warning: /usr/local/sqoop/./hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /usr/local/sqoop/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
Warning: /usr/local/sqoop/./zookeeper does not exist! Accumulo imports will fail.
Please set $ZOOKEEPER_HOME to the root of your Zookeeper installation.
2017-11-18 19:44:29,360 INFO [main] sqoop.Sqoop: Running Sqoop version: 1.4.5
Enter password:
2017-11-18 19:44:31,593 INFO [main] manager.MySQLManager: Preparing to use a MySQL streaming resultset.
2017-11-18 19:44:31,593 INFO [main] tool.CodeGenTool: Beginning code generation
2017-11-18 19:44:32,094 INFO [main] manager.SqlManager: Executing SQL statement: SELECT t.* FROM `employee` AS t LIMIT 1
2017-11-18 19:44:32,202 INFO [main] manager.SqlManager: Executing SQL statement: SELECT t.* FROM `employee` AS t LIMIT 1
2017-11-18 19:44:32,236 INFO [main] orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/local/hadoop-2.6.0
Note: /tmp/sqoop-acadgild/compile/1a4c91a86480b6264ed716c501b7d3ff/employee.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
2017-11-18 19:44:39,447 INFO [main] orm.CompilationManager: Writing jar file: /tmp/sqoop-acadgild/compile/1a4c91a86480b6264ed716c501b7d3ff/employee.jar
```

- Check the 'sqoopout' directory in HDFS

```
hadoop fs -ls /sqoopout
hadoop fs -cat /sqoopout/part-m-00000
```

```
[acadgild@localhost bin]$ hadoop fs -ls /sqoopout
17/11/18 19:45:40 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform: java.lang.UnsupportedClassVersionError: org.apache.hadoop.util.NativeCodeLoader:1.0 incompatible with JVM version 1.8.0_102
Found 2 items
-rw-r--r-- 1 acadgild supergroup 0 2017-11-18 19:45 /sqoopout/_SUCCESS
-rw-r--r-- 1 acadgild supergroup 167 2017-11-18 19:45 /sqoopout/part-m-00000
[acadgild@localhost bin]$ hadoop fs -cat /sqoopout/part-m-00000
17/11/18 19:46:03 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform: java.lang.UnsupportedClassVersionError: org.apache.hadoop.util.NativeCodeLoader:1.0 incompatible with JVM version 1.8.0_102
1,Mohan,25,Big Data & Hadoop,30000
2,Ramu,27,AI,50000
3,Ravi,30,Java,60000
4,Akshith,22,Automation,35000
5,Shyam,35,C,40000
6,Priya,28,.Net,50000
7,Madhu,27,DBA,70000
[acadgild@localhost bin]$
```

Export: Transfer of data from HDFS to Mysql

- **Delete data from table 'employee' in Mysql**

use db1;

delete from employee;

```
mysql> select * from employee;
+----+-----+-----+-----+-----+
| id | name  | age | skill                | salary |
+----+-----+-----+-----+-----+
| 1  | Mohan | 25  | Big Data & Hadoop    | 30000  |
| 2  | Ramu  | 27  | AI                   | 50000  |
| 3  | Ravi  | 30  | Java                 | 60000  |
| 4  | Akshith | 22  | Automation           | 35000  |
| 5  | Shyam | 35  | C                    | 40000  |
| 6  | Priya | 28  | .Net                 | 50000  |
| 7  | Madhu | 27  | DBA                  | 70000  |
+----+-----+-----+-----+-----+
7 rows in set (0.00 sec)

mysql> delete from employee;
Query OK, 7 rows affected (0.00 sec)

mysql>
```

- **Export data using below query from directory '/sqoopout' to table 'employee' in Mysql**

sqoop export --connect jdbc:mysql://localhost/assignment11 --username 'root' -P --table 'employee' --export-dir '/sqoopout' --input-fields-terminated-by ';' -m 1 -c columns id,name,age,salary

```
[acadgild@localhost ~]$ sqoop export --connect jdbc:mysql://localhost/assignment11 --username 'root' -P --table 'employee' --export-dir '/sqoopout' --input-fields-terminated-by ';' -m 1 --columns id,name,age,salary
Warning: /usr/local/sqoop/../hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /usr/local/sqoop/../accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
Warning: /usr/local/sqoop/../zookeeper does not exist! Accumulo imports will fail.
Please set $ZOOKEEPER_HOME to the root of your Zookeeper installation.
2017-11-18 19:58:37,495 INFO [main] sqoop.Sqoop: Running Sqoop version: 1.4.5
Enter password:
2017-11-18 19:58:38,634 INFO [main] manager.MySQLManager: Preparing to use a MySQL streaming resultset.
2017-11-18 19:58:38,634 INFO [main] tool.CodeGenTool: Beginning code generation
2017-11-18 19:58:39,240 INFO [main] manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'employee' AS t LIMIT 1
2017-11-18 19:58:39,304 INFO [main] manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'employee' AS t LIMIT 1
2017-11-18 19:58:39,343 INFO [main] orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/local/hadoop-2.6.0
Note: /tmp/sqoop-acadgild/compile/2583e7b5b6cef8a745e0f1f32a0f6aec/employee.java uses or overrides a deprecated API.
```

- **Check table 'employee' in Mysql for data.**

mysql -u root

use db1;

```
select * from employee;
```

```
mysql> select * from employee;
+----+-----+-----+-----+-----+
| id | name  | age  | skill                | salary |
+----+-----+-----+-----+-----+
| 1  | Mohan | 25   | Big Data & Hadoop    | 30000  |
| 2  | Ramu  | 27   | AI                   | 50000  |
| 3  | Ravi  | 30   | Java                 | 60000  |
| 4  | Akshith | 22  | Automation           | 35000  |
| 5  | Shyam | 35   | C                    | 40000  |
| 6  | Priya | 28   | .Net                 | 50000  |
| 7  | Madhu | 27   | DBA                  | 70000  |
+----+-----+-----+-----+-----+
7 rows in set (0.00 sec)

mysql> █
```

Transfer data between Mysql and Hive (Import and Export only selected columns) using Sqoop.

Import: Transfer of data from Mysql to hive.

- **Create table in Mysql**

```
create database assignment11;
use assignment11;
create table employee
(id int, name varchar(20),
age int, skill varchar(20),
salary integer(10));
```

```
mysql> show databases;
+-----+
| Database |
+-----+
| information_schema |
| b1 |
| db1 |
| metastore |
| mysql |
| test |
+-----+
6 rows in set (0.00 sec)

mysql> create database assignment11;
Query OK, 1 row affected (0.00 sec)

mysql> use assignment11;
Database changed
mysql> create table employee
-> (id int, name varchar(20),
-> age int, skill varchar(20),
-> salary integer(10));
Query OK, 0 rows affected (0.01 sec)

mysql> show databases;
+-----+
| Database |
+-----+
| information_schema |
| assignment11 |
| b1 |
| db1 |
| metastore |
| mysql |
| test |
+-----+
7 rows in set (0.00 sec)
```

- **Insert Data to created table: 'employee'**

```
insert into employee values(1, 'Mohan',25,'Big Data & Hadoop',30000);
insert into employee values(2, 'Ramu',27,'AI',50000);
insert into employee values(3, 'Ravi',30,'Java',60000);
insert into employee values(4, 'Akshith',22,'Automation',35000);
insert into employee values(5, 'Shyam',35,'C',40000);
insert into employee values(6, 'Priya',28,'.Net',50000);
insert into employee values(7, 'Madhu',27,'DBA',70000);
```

**Grant all privileges to root user and commit the data.**

```
grant all on *.* to 'root'@'localhost' with grant option;
flush privileges;
commit;
exit;
```

```
mysql> select * from employee;
Empty set (0.00 sec)

mysql> insert into employee values(1, 'Mohan',25,'Big Data & Hadoop',30000);
Query OK, 1 row affected (0.00 sec)

mysql> insert into employee values(2, 'Ramu',27,'AI',50000);
Query OK, 1 row affected (0.00 sec)

mysql> insert into employee values(3, 'Ravi',30,'Java',60000);
Query OK, 1 row affected (0.00 sec)

mysql> insert into employee values(4, 'Akshith',22,'Automation',35000);
Query OK, 1 row affected (0.00 sec)

mysql> insert into employee values(5, 'Shyam',35,'C',40000);
Query OK, 1 row affected (0.00 sec)

mysql> insert into employee values(6, 'Priya',28,'.Net',50000);
Query OK, 1 row affected (0.00 sec)

mysql> insert into employee values(7, 'Madhu',27,'DBA',70000);
Query OK, 1 row affected (0.00 sec)

mysql> select * from employee;
+----+-----+-----+-----+-----+
| id  | name  | age  | skill                | salary |
+----+-----+-----+-----+-----+
| 1   | Mohan | 25   | Big Data & Hadoop    | 30000  |
| 2   | Ramu  | 27   | AI                   | 50000  |
| 3   | Ravi  | 30   | Java                 | 60000  |
| 4   | Akshith | 22  | Automation           | 35000  |
| 5   | Shyam | 35   | C                    | 40000  |
| 6   | Priya | 28   | .Net                 | 50000  |
| 7   | Madhu | 27   | DBA                  | 70000  |
+----+-----+-----+-----+-----+
7 rows in set (0.00 sec)
```

- **Use below query to transfer data from Mysql table 'employee' to Hive**  
sqoop import \  
--connect jdbc:mysql://localhost/assignment11 \  
--username 'root' --table 'employee' --target-dir '/sqoopout' \  
--hive-import \  
-m 1;



```

[root@sandbox ~]# sqoop import \
> --connect jdbc:mysql://localhost/assignment11 \
> --username 'root' --table 'employee' --target-dir '/sqoopout' \
> --hive-import \
> -m 1;
Warning: /usr/hdp/2.2.0.0-2041/accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
17/11/28 02:00:54 INFO sqoop.Sqoop: Running Sqoop version: 1.4.5.2.2.0.0-2041
17/11/28 02:00:54 INFO tool.BaseSqoopTool: Using Hive-specific delimiters for output. You can override
17/11/28 02:00:54 INFO tool.BaseSqoopTool: delimiters with --fields-terminated-by, etc.
17/11/28 02:00:56 INFO manager.SqlManager: Using default fetchSize of 1000
17/11/28 02:00:56 INFO tool.CodeGenTool: Beginning code generation
17/11/28 02:00:57 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'employee' AS t LIMIT 1
17/11/28 02:00:57 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'employee' AS t LIMIT 1
17/11/28 02:00:57 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/hdp/2.2.0.0-2041/hadoop-mapreduce
Note: /tmp/sqoop-root/compile/c53e0d790eb2751fbl58f75ccb8c2ee/employee.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
17/11/28 02:01:09 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-root/compile/c53e0d790eb2751fbl58f75ccb8c2ee/employee.jar
17/11/28 02:01:10 WARN manager.MySQLManager: It looks like you are importing from mysql.
17/11/28 02:01:10 WARN manager.MySQLManager: This transfer can be faster! Use the --direct
17/11/28 02:01:10 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.
17/11/28 02:01:10 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)
17/11/28 02:01:10 INFO mapreduce.ImportJobBase: Beginning import of employee
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/hdp/2.2.0.0-2041/hadoop/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/hdp/2.2.0.0-2041/zookeeper/lib/slf4j-log4j12-1.6.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/hdp/2.2.0.0-2041/hive/lib/hive-jdbc-0.14.0.2.2.0.0-2041-standalone.jar!/org/slf4j/impl/StaticLogge
rBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
17/11/28 02:01:26 INFO impl.TimelineClientImpl: Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/
17/11/28 02:01:27 INFO client.RMPProxy: Connecting to ResourceManager at sandbox.hortonworks.com/10.0.2.15:8050
17/11/28 02:02:36 INFO db.DBInputFormat: Using read committed transaction isolation
17/11/28 02:02:40 INFO mapreduce.JobSubmitter: number of splits:1
17/11/28 02:02:46 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1511833512519_0001
17/11/28 02:03:21 INFO impl.YarnClientImpl: Submitted application application_1511833512519_0001
17/11/28 02:03:47 INFO mapreduce.Job: The url to track the job: http://sandbox.hortonworks.com:8088/proxy/application_1511833512519_0001/
17/11/28 02:03:47 INFO mapreduce.Job: Running job: job_1511833512519_0001
17/11/28 02:02:40 INFO mapreduce.JobSubmitter: number of splits:1
17/11/28 02:02:46 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1511833512519_0001
17/11/28 02:03:21 INFO impl.YarnClientImpl: Submitted application application_1511833512519_0001
17/11/28 02:03:47 INFO mapreduce.Job: The url to track the job: http://sandbox.hortonworks.com:8088/proxy/application_1511833512519_0001/
17/11/28 02:03:47 INFO mapreduce.Job: Running job: job_1511833512519_0001
17/11/28 02:07:51 INFO mapreduce.Job: Job job_1511833512519_0001 running in uber mode : false
17/11/28 02:07:51 INFO mapreduce.Job: map 0% reduce 0%
17/11/28 02:11:12 INFO mapreduce.Job: map 100% reduce 0%
17/11/28 02:12:53 INFO mapreduce.Job: Job job_1511833512519_0001 completed successfully
17/11/28 02:14:50 INFO mapreduce.Job: Counters: 30
  File System Counters
    FILE: Number of bytes read=0
    FILE: Number of bytes written=123779
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=87
    HDFS: Number of bytes written=167
    HDFS: Number of read operations=4
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Other local map tasks=1
    Total time spent by all maps in occupied slots (ms)=204235
    Total time spent by all reduces in occupied slots (ms)=0
    Spilled Records=0
    Failed Shuffles=0
    Merged Map outputs=0
    GC time elapsed (ms)=139
    CPU time spent (ms)=1990
    Physical memory (bytes) snapshot=104423424
    Virtual memory (bytes) snapshot=785715200
    Total committed heap usage (bytes)=39321600
  File Input Format Counters
    Bytes Read=0
  File Output Format Counters
    Bytes Written=167
17/11/28 02:14:52 INFO ipc.Client: Retrying connect to server: sandbox.hortonworks.com/10.0.2.15:36711. Already tried 0 time(s); retry po
lICY is RetryUpToMaximumCountWithFixedSleep(maxRetries=3, sleepTime=1000 MILLISECONDS)
17/11/28 02:14:53 INFO ipc.Client: Retrying connect to server: sandbox.hortonworks.com/10.0.2.15:36711. Already tried 1 time(s); retry po
lICY is RetryUpToMaximumCountWithFixedSleep(maxRetries=3, sleepTime=1000 MILLISECONDS)
17/11/28 02:14:54 INFO ipc.Client: Retrying connect to server: sandbox.hortonworks.com/10.0.2.15:36711. Already tried 2 time(s); retry po
lICY is RetryUpToMaximumCountWithFixedSleep(maxRetries=3, sleepTime=1000 MILLISECONDS)
17/11/28 02:15:05 INFO mapred.ClientServiceDelegate: Application state is completed. FinalApplicationStatus=SUCCEEDED. Redirecting to job
history server
17/11/28 02:16:28 INFO mapreduce.ImportJobBase: Transferred 167 bytes in 905.9278 seconds (0.1843 bytes/sec)
17/11/28 02:16:28 INFO mapreduce.ImportJobBase: Retrieved 7 records.
17/11/28 02:16:34 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'employee' AS t LIMIT 1
17/11/28 02:17:26 INFO hive.HiveImport: Loading uploaded data into Hive
17/11/28 02:17:56 WARN conf.HiveConf: HiveConf of name hive.optimize.mapjoin.mapreduce does not exist
17/11/28 02:17:56 WARN conf.HiveConf: HiveConf of name hive.heapsize does not exist
17/11/28 02:17:56 WARN conf.HiveConf: HiveConf of name hive.server2.enable.imperpersonation does not exist
17/11/28 02:17:56 WARN conf.HiveConf: HiveConf of name hive.auto.convert.sortmerge.join.noconditionaltask does not exist
Logging initialized using configuration in jar:file:/usr/hdp/2.2.0.0-2041/hive/lib/hive-common-0.14.0.2.2.0.0-2041.jar!/hive-log4j.proper
ties
OK
Time taken: 232.608 seconds
Loading data to table default.employee
Table default.employee stats: [numFiles=1, totalSize=167]
OK
Time taken: 71.856 seconds
[root@sandbox ~]#

```

- Check out hive:



```

hive> show databases;
OK
default
xademo
Time taken: 13.316 seconds, Fetched: 2 row(s)
hive> use default;
OK
Time taken: 7.769 seconds
hive> show tables;
OK
employee
sample_07
sample_08
Time taken: 2.815 seconds, Fetched: 3 row(s)
hive> select * from employee;
OK
1      Mohan    25      Big Data & Hadoop      30000
2      Ramu     27      AI                    50000
3      Ravi     30      Java                  60000
4      Akshith  22      Automation            35000
5      Shyam    35      C                     40000
6      Priya    28      .Net                  50000
7      Madhu    27      DBA                   70000
Time taken: 10.728 seconds, Fetched: 7 row(s)
hive> █

```

Export: Transfer of data from hive to Mysql.

- **Delete data from table 'employee' in Mysql**

use db1;

delete from employee;

```

mysql> use assignment11;
Database changed
mysql> show tables;
+-----+
| Tables_in_assignment11 |
+-----+
| employee                |
+-----+
1 row in set (0.00 sec)

mysql> select * from employee;
+-----+-----+-----+-----+-----+
| id  | name  | age  | skill                | salary |
+-----+-----+-----+-----+-----+
| 1   | Mohan | 25   | Big Data & Hadoop    | 30000  |
| 2   | Ramu  | 27   | AI                   | 50000  |
| 3   | Ravi  | 30   | Java                 | 60000  |
| 4   | Akshith | 22  | Automation           | 35000  |
| 5   | Shyam | 35   | C                    | 40000  |
| 6   | Priya | 28   | .Net                 | 50000  |
| 7   | Madhu | 27   | DBA                  | 70000  |
| 8   | Suraj | 28   | Team Lead            | 80000  |
| 9   | Ganesh | 30  | Manager              | 100000 |
+-----+-----+-----+-----+-----+
9 rows in set (0.03 sec)

```

```
mysql> delete from employee;
Query OK, 9 rows affected (0.03 sec)

mysql> select * from employee;
Empty set (0.01 sec)

mysql>
```

- Export data using below query from hive to table 'employee' in Mysql

```
sqoop export --connect jdbc:mysql://localhost/assignment11 -username root --columns id,name,skill
--table employee --export-dir
hdfs://sandbox.hortonworks.com:8020/apps/hive/warehouse/employee --input-fields-terminated-
by '\u0001' -m 1
```

```
[root@sandbox ~]# sqoop export --connect jdbc:mysql://localhost/assignment11 -username root --columns id,name,skill --table employee --ex
port-dir hdfs://sandbox.hortonworks.com:8020/apps/hive/warehouse/employee --input-fields-terminated-by '\u0001' -m 1
Warning: /usr/hdp/2.2.0.0-2041/accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
17/12/04 16:26:54 INFO sqoop.Sqoop: Running Sqoop version: 1.4.5.2.2.0-2041
17/12/04 16:26:54 WARN sqoop.SqoopOptions: Character argument '\u0001' has multiple characters; only the first will be used.
17/12/04 16:26:55 INFO manager.SqlManager: Using default fetchSize of 1000
17/12/04 16:26:55 INFO tool.CodeGenTool: Beginning code generation
17/12/04 16:26:55 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'employee' AS t LIMIT 1
17/12/04 16:26:56 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'employee' AS t LIMIT 1
17/12/04 16:26:56 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/hdp/2.2.0.0-2041/hadoop-mapreduce
Note: /tmp/sqoop-root/compile/ad179be3475af08e9dd4cce8c9cb0e2d/employee.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
17/12/04 16:27:04 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-root/compile/ad179be3475af08e9dd4cce8c9cb0e2d/employee.jar
17/12/04 16:27:04 INFO mapreduce.ExportJobBase: Beginning export of employee
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/hdp/2.2.0.0-2041/hadoop/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/hdp/2.2.0.0-2041/zookeeper/lib/slf4j-log4j12-1.6.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/hdp/2.2.0.0-2041/hive/lib/hive-jdbc-0.14.0.2.2.0-2041-standalone.jar!/org/slf4j/impl/StaticLogge
rBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
17/12/04 16:27:21 INFO impl.TimelineClientImpl: Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/
17/12/04 16:27:22 INFO client.RMProxy: Connecting to ResourceManager at sandbox.hortonworks.com:10.0.2.15:8050
17/12/04 16:27:34 INFO input.FileInputFormat: Total input paths to process : 2
17/12/04 16:27:34 INFO input.FileInputFormat: Total input paths to process : 2
17/12/04 16:27:35 INFO mapreduce.JobSubmitter: number of splits:1
17/12/04 16:27:36 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1512403504625_0001
17/12/04 16:27:43 INFO impl.YarnClientImpl: Submitted application application_1512403504625_0001
17/12/04 16:27:52 INFO mapreduce.Job: The url to track the job: http://sandbox.hortonworks.com:8088/proxy/application_1512403504625_0001/
```

- Check table 'employee' in Mysql for data.

```
mysql -u root
use db1;
select * from employee;
```

```
mysql> select * from employee;
+----+-----+-----+-----+-----+
| id | name  | age  | skill                | salary |
+----+-----+-----+-----+-----+
| 1  | Mohan | 25   | Big Data & Hadoop    | 30000  |
| 2  | Ramu  | 27   | AI                   | 50000  |
| 3  | Ravi  | 30   | Java                 | 60000  |
| 4  | Akshith | 22  | Automation           | 35000  |
| 5  | Shyam | 35   | C                    | 40000  |
| 6  | Priya | 28   | .Net                 | 50000  |
| 7  | Madhu | 27   | DBA                  | 70000  |
| 8  | Suraj | 28   | Team Lead            | 80000  |
| 9  | Ganesh | 30  | Manager              | 100000 |
+----+-----+-----+-----+-----+
9 rows in set (0.04 sec)
```